

Database Applications


Lecture 5: Views, Triggers, Stored Procedures

Santha Sumanasekara
August 2019



1

1



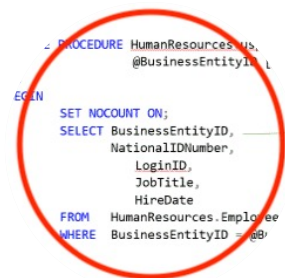
Views

Views
Materialized Views



Triggers

Before Triggers
After Triggers
INSTEAD OF Triggers



Stored Procedures

PL/SQL
Anonymous Blocks
Stored Procedures



2

2

“Simple” Views– why?

- Views are kind of virtual tables, allow users to do the following:
 - Structure data in a way that users find natural or intuitive.
 - Restrict access to the data such that a user can only see limited data instead of complete table.
 - Summarise data from various tables which can be used to generate reports.
- Views are defined using an SQL statement. So, you may consider a view as a stored query, assigned with a name (and also a virtual schema derived from the underlying SQL statement).



3

3

Views – Create a view

- On our movie rental database, we can summarise data, say for each film, we can count the number copies in the inventory, and store along with the movie title:

SQL | Fetched 100 rows in 0.048 seconds

| FILM_ID | TITLE | Copy Count |
|---------|-----------------------|------------|
| 39 | 128 CATCH AMISTAD | 0 |
| 40 | 642 ORDER BETRAYED | 0 |
| 41 | 712 RAIDERS ANTITRUST | 0 |
| 42 | 192 CROSSING DIVORCE | 0 |
| 43 | 721 REDS POCUS | 2 |
| 44 | 566 MAUDE MOD | 2 |
| 45 | 362 GLORY TRACY | 2 |
| 46 | 885 TEXAS WATCH | 2 |
| 47 | 910 TREATMENT JEKYLL | 2 |
| 48 | 799 SIMON NORTH | 2 |



4

4

Views – Create a view

```
SELECT f.film_id, title, COUNT(i.film_id) AS "Copy Count"
FROM film f LEFT OUTER JOIN inventory i ON f.film_id = i.film_id
GROUP BY f.film_id, title
ORDER BY COUNT(i.film_id);
```

SQL | Fetched 100 rows in 0.048 seconds

| FILM_ID | TITLE | Copy Count |
|---------|-----------------------|------------|
| 39 | 128 CATCH AMISTAD | 0 |
| 40 | 642 ORDER BETRAYED | 0 |
| 41 | 712 RAIDERS ANTITRUST | 0 |
| 42 | 192 CROSSING DIVORCE | 0 |
| 43 | 721 REDS POCUS | 2 |
| 44 | 566 MAUDE MOD | 2 |
| 45 | 362 GLORY TRACY | 2 |
| 46 | 885 TEXAS WATCH | 2 |
| 47 | 910 TREATMENT JEKYLL | 2 |
| 48 | 799 SIMON NORTH | 2 |

Views – Create a view

```
SELECT f.film_id, title, COUNT(i.film_id) AS "Copy Count"
FROM film f LEFT OUTER JOIN inventory i ON f.film_id = i.film_id
GROUP BY f.film_id, title
ORDER BY COUNT(i.film_id);
```

- Use the above SQL query as the basis of the view definition.

```
CREATE VIEW film_copy (film_id, title, copy_count)
AS
```

```
SELECT f.film_id, title, COUNT(i.film_id) AS "Copy Count"
FROM film f LEFT OUTER JOIN inventory i ON f.film_id = i.film_id
GROUP BY f.film_id, title
ORDER BY COUNT(i.film_id);
```

Views – Using a view

- A view can be use exactly the same way as a table.

```
SELECT *
FROM film_copy;
```

- When you run this query, data are extracted from base tables.
- Data are not physically stored on the view.
- On the view definition is stored.

Film_copy is the view we created in the previous step.



7

7

Views – Using a view

| Worksheet Query Builder | | | | | |
|---|-------------|----------|-------------|------|----|
| select * from film_copy | | | | | |
| Script Output x Query Result x Explain Plan x | | | | | |
| SQL 0.147 seconds | | | | | |
| OPERATION | OBJECT_NAME | OPTIONS | CARDINALITY | COST | |
| SELECT STATEMENT | | | | 4623 | 34 |
| VIEW | FILM_COPY | | | 4623 | 34 |
| SORT | | ORDER BY | | 4623 | 34 |
| HASH | | GROUP BY | | 4623 | 34 |
| HASH JOIN | | OUTER | | 4623 | 29 |
| Access Predicates | | | | | |
| F.FILM_ID=I.FILM_ID(+) | | | | | |
| TABLE ACCESS | FILM | FULL | | 1000 | 19 |
| TABLE ACCESS | INVENTORY | FULL | | 4581 | 9 |



8

8

Views – Important points

- A view is not a data source, it is only a definition;
- Data always reside base tables;
- As seen on Execution plan, no performance gains can be achieved;
- Only an abstract view of one or more data sources.



A good way to hide complex table joins and restricted columns.

9

9

Materiali[s|z]ed Views

Mother of all query optimisations!



10

Materiali[s|z]ed Views

- A "materialised view" is different to a "view" – they store data!
- A MV is defined using an SQL query. At the point of CREATE MATERIALIZED VIEW, the base query is run, data extracted from base tables and stored in the MV.
- So, they take disk space
- Also, integrity of data is a major concern.



A good way to hide complex table joins and restricted columns.

11

11

Materiali[s|z]ed Views -- Example

- Consider a query to aggregation of on-time flight data.
- How many flights operated by each aircraft (identified by the tail no)?

```
SELECT tailnum, COUNT(*)
  FROM ontime
 GROUP BY talinum;
```

- Considering about 7 million rows and full-table scan as the only viable access method, this will take time!



12

12

Query Editor

Query History

```

1 SELECT tailnum,COUNT(*)
2 FROM ontime
3 GROUP BY tailnum
4

```

Explain

Data Output

Notifications

Messages

public.ontime

Aggregate

Sort

Gather Merge

Aggregate

13

13

Query Editor

Query History

```

1 EXPLAIN ANALYZE SELECT tailnum,COUNT(*)
2 FROM ontime
3 GROUP BY tailnum
4

```

Explain

Data Output

Messages

QUERY PLAN

text

1

Finalize GroupAggregate (cost=168689.71..169949.87 rows=4974 width=14) (actual time=8599.584..8610.021...

2

Group Key: tailnum

3

-> Gather Merge (cost=168689.71..169850.39 rows=9948 width=14) (actual time=8599.565..8606.346 rows=...

4

Workers Planned: 2

5

Workers Launched: 2

6

-> Sort (cost=167689.68..167702.12 rows=4974 width=14) (actual time=8593.925..8594.699 rows=5358 l...

7

Sort Key: tailnum

8

Sort Method: quicksort Memory: 444kB

9

Worker 0: Sort Method: quicksort Memory: 443kB

10

Worker 1: Sort Method: quicksort Memory: 444kB

11

-> Partial HashAggregate (cost=167334.54..167384.27 rows=4974 width=14) (actual time=8588.397.....

12

Group Key: tailnum

13

-> Parallel Seq Scan on ontime (cost=0.00..152731.02 rows=2920702 width=6) (actual time=0.128...

14

Planning Time: 1.065 ms

15

Execution Time: 8611.499 ms

14

14

Materiali[s|z]ed Views -- Example

- Let's suppose we already have a pre-built materialized view – `ontime_mv`
- `ontime_mv` was built using the same SQL query.

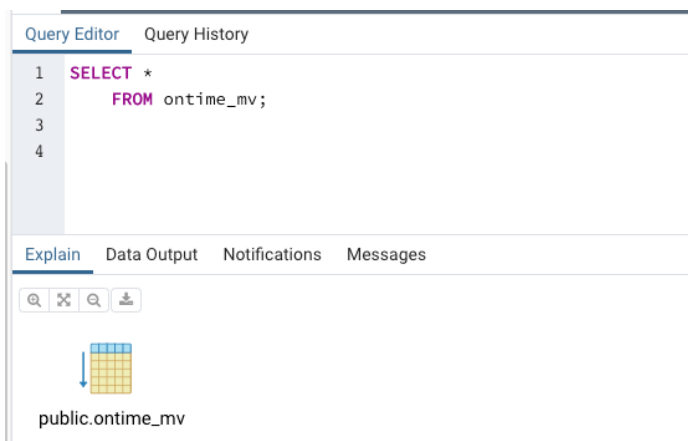
```
SELECT *  
FROM ontime_mv;
```

- Now look at query performance.



15

15



16

16

Query Editor Query History

```

1  EXPLAIN ANALYZE SELECT *
2    FROM ontime_mv;
3
4

```

8000 seconds vs. 1 second!

Explain **Data Output** Notifications Messages

| | QUERY PLAN |
|---|---|
| 1 | Seq Scan on ontime_mv (cost=0.00..83.74 rows=5374 width=14) (actual time=0.018..... |
| 2 | Planning Time: 0.026 ms |
| 3 | Execution Time: 1.035 ms |

UNIVERSITY

17

17

Materialized Views -- Example

- So enticing -- how `ontime_mv` was built?
- `ontime_mv` was built using the same SQL query.

```

CREATE MATERIALIZED VIEW ontime_mv
AS
  SELECT tailnum, COUNT(*)
     FROM ontime
    GROUP BY talinum;

```

18

MVs in Data Warehousing

- In data warehouses, materialized views are used to compute and store aggregated data such as sums and averages.
- Materialized views in these environments are typically referred to as summaries because they store summarised data.
- The query optimiser can use materialized views to improve query performance by automatically recognising when a materialized view can and should be used to satisfy a request.



19

19

Refreshing MVs

- A methodical way to refresh data in MVs is very important.
- Whenever, INSERT, UPDATE or DELETE data (any DML operation), data in related MVs become inconsistent with base tables.
- What can we do to ensure they are consistent?
- Refresh MVs.
- How? When?



20

20

Refreshing MVs -- How?

- There are two established ways of refreshing MVs.

- Fast Refresh

Apply changes made after last refresh. A Materialized View Log is maintained to keep track of changes.

- Full Refresh

Re-run the whole query and rebuild from scratch.



21

21

Refreshing MVs -- When?

- There are two established ways of refreshing MVs.

- On Demand

Changes are made when you ask the database to do so. Between the DML and the refresh point, MV still contain old data

- On Commit

When DML is committed MV is refreshed.



22

22

Refreshing MVs – How and When?

- Not all combinations are supported.

```
CREATE MATERIALIZED VIEW ontime_mv  
REFRESH FAST ON COMMIT  
AS  
    SELECT tailnum, COUNT(*)  
        FROM ontime  
        GROUP BY talinum;
```

Refreshing MVs – How and When?

- Not all combinations are supported.

```
CREATE MATERIALIZED VIEW ontime_mv  
REFRESH FULL ON DEMAND  
AS  
    SELECT tailnum, COUNT(*)  
        FROM ontime  
        GROUP BY talinum;
```

Refreshing MVs – How and When?

- Not all combinations are supported.

| When/ How | Fast | Complete |
|--------------|------|----------|
| On Demand | Yes | Yes |
| On Commit | Yes | May be |
| On Statement | Yes | No |

Subject to many restrictions.

Only on new versions of Oracle. (12.2 and above)

Compare MVs vs Views

| | View | Materialized View |
|----------------|------------------|-------------------------------|
| Define | CREATE VIEW v AS | CREATE MATERIALIZED VIEW v AS |
| Store | Definition only | Definition and data |
| Speed | No gain | Very fast |
| Data Integrity | No issues | Must be carefully managed |

Triggers

Sorry, not about “Tigger” 😊!



27

RMIT Classification: Trusted

Triggers

- Triggers are operations that are automatically performed when a specified database event occurs.
- A trigger event can be a changing data action -- INSERT, UPDATE, or DELETE.
- A trigger action can be another action (SELECT, INSERT, UPDATE, DELETE) action.

28

RMIT Classification: Trusted

Triggers -- Examples

| Trigger Event | Trigger Action |
|--|--|
| updating a table | add a log record to another table (say, <i>action_log</i>) to store timestamp and old and new values. |
| inserting a new record into Sales table | updating Inventory table. |
| a user view an item on an E-Commerce website | update a <i>User_Recommender</i> table for targeted marketing. |
| User wants to update a view | trigger will update the underlying tables Instead! |

29

RMIT Classification: Trusted

Triggers – Three types

- BEFORE triggers
- AFTER Triggers
- INSTEAD OF triggers

30

RMIT Classification: Trusted

BEFORE Triggers

- BEFORE triggers execute trigger actions before the trigger event is executed.
- An example is adding a log record **before** doing the original action.
- E.g. When updating a table add a log record to another table (say, action_log) to store timestamp and old and new values.



31

RMIT Classification: Trusted

BEFORE Triggers – An example

- Film table has an attribute 'rental_rate'.
- Rental rate is variable, and time to time we update them.
- Whenever we update the current rental rate of a film, we must keep a log record to store the old rental rate and new rental rate and the timestamp
- This helps us to keep track of historical changes of rental rates.



32

RMIT Classification: Trusted

BEFORE Triggers – An example

```
CREATE OR REPLACE TRIGGER update_film_log
  BEFORE UPDATE OF rental_rate ON film
FOR EACH ROW
BEGIN
  INSERT INTO film_log VALUES
    (:NEW.film_id, :OLD.rental_rate, :NEW.rental_rate,
    sysdate);
END;
```

Trigger Event

Trigger Action



33

RMIT Classification: Trusted

BEFORE Triggers – An example

```
CREATE OR REPLACE TRIGGER update_film_log
  BEFORE UPDATE OF rental_rate ON film
FOR EACH ROW
BEGIN
  INSERT INTO film_log VALUES
    (:NEW.film_id, :OLD.rental_rate, :NEW.rental_rate,
    sysdate);
END;
```

:NEW.film_id refers
to the film that is
being updated

:OLD.rental_rate refers
to the current rental rate
value of the
corresponding row in
the film table.



34

RMIT Classification: Trusted

BEFORE Triggers – An example

```
UPDATE film SET
rental_rate = 3.99
WHERE LOWER(title) = 'metropolis coma'
```

Diagram illustrating the execution of an UPDATE statement and the resulting data changes in the **film** and **film_Log** tables.

film table (Before):

| FILM_ID | TITLE | RENTAL_RATE |
|---------|---------------------|-------------|
| 1 | 572 METROPOLIS COMA | 2.99 |

film table (After):

| FILM_ID | TITLE | RENTAL_RATE |
|---------|---------------------|-------------|
| 1 | 572 METROPOLIS COMA | 3.99 |

film_Log table (After):

| FILM_ID | OLD_RENTAL_RATE | NEW_RENTAL_RATE | CHANGE_DATE |
|---------|-----------------|-----------------|--------------------------------------|
| 1 | 572 | 2.99 | 3.99 22/AUG/19 11:48:26.000000000 AM |

This UPDATE statement fires the "update trigger".

35

RMIT Classification: Trusted

AFTER Triggers

- AFTER triggers execute trigger actions after the trigger event is executed.
- An example is display before and after values **after** doing and update action.
- Operation of the AFTER trigger is very similar to BEFORE trigger, only the timing is different.

36

RMIT Classification: Trusted

INSTEAD OF Triggers

- INSTEAD OF triggers are mainly used for updating VIEWS.
- When a user asks for updating a view, normally it is not permitted. However, you can use a trigger to update the underlying tables instead.



37

RMIT Classification: Trusted

INSTEAD OF Triggers – An example

- A user wants to update the name of a continent.

```
UPDATE film_copy
SET title = 'METROPOLIS DEATH'
WHERE TITLE = 'METROPOLIS COMA';
```

```
Error starting at line : 1 in command -
UPDATE film_copy
  SET title = 'METROPOLIS DEATH'
  WHERE TITLE = 'METROPOLIS COMA'
Error at Command Line : 1 Column : 8
Error report -
SQL Error: ORA-01732: data manipulation operation not legal on this view
01732. 00000 - "data manipulation operation not legal on this view"
*Cause:
*Action:
```

Because film_copy is a view.



38

RMIT Classification: Trusted

INSTEAD OF Triggers – An example

- Define an INSTEAD OF trigger to do the job.

```
CREATE OR REPLACE trigger view_update
  INSTEAD OF UPDATE ON film_copy
BEGIN
  UPDATE film
    SET title = :NEW.title
    WHERE title = :OLD.title;
END;
```

Trigger action – only
the action get
executed.

Trigger event – not
get executed at all.

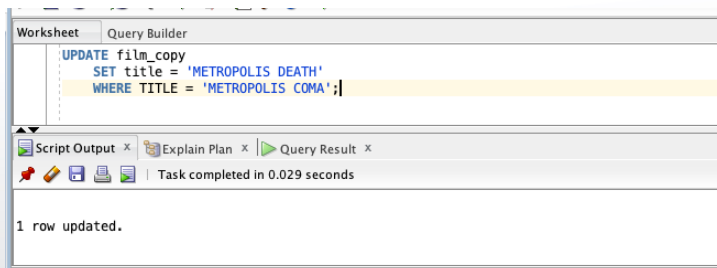


39

RMIT Classification: Trusted

INSTEAD OF Triggers – An example

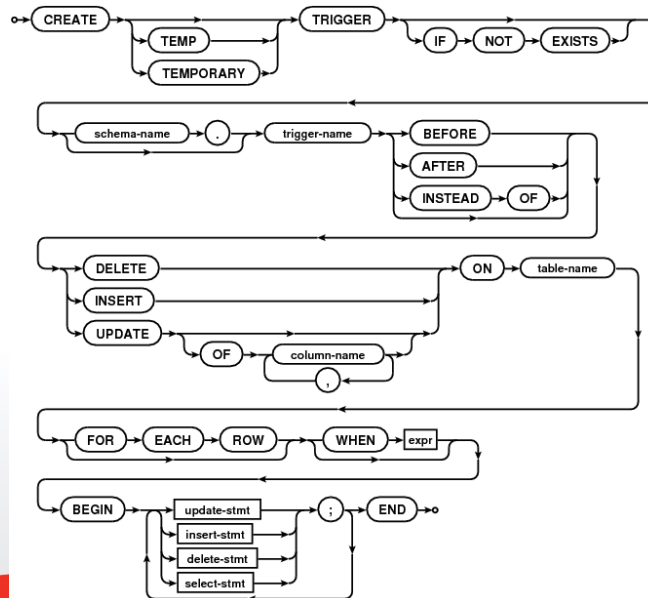
```
CREATE OR REPLACE trigger view_update
  INSTEAD OF UPDATE ON film_copy
BEGIN
  UPDATE film
    SET title = :NEW.title
    WHERE title = :OLD.title;
END
```



40

RMIT Classification: Trusted

CREATE TRIGGER syntax



41

RMIT Classification: Trusted

About use of OLD and NEW

- When executing a trigger, the database engine allows you to refer old and new rows the trigger event is dealing with.
- They can be used within trigger action.
- Depend on the trigger event, we may use either of them or both.

| Trigger Event | Available references |
|---------------|----------------------|
| INSERT | :NEW |
| UPDATE | :NEW, :OLD |
| DELETE | :OLD |

42

Stored Procedures

Extending the capabilities of SQL!



43

Stored Procedures

- SQL is a declarative query language.
- Using a declarative database query language may also result in better code than what can be created manually.
- Declarative query languages are also easier to use as they simply focus on what must be retrieved and do so quickly. However, declarative languages have their own trade-offs.
- Users have little to no control over how inputs are dealt with.
- If the user wants to use a functionality that the query language doesn't support, they are often at a loss.



44

44

Stored Procedures

Extending the capabilities of SQL!



45

Stored Procedures

- PL/SQL is Oracle's procedural language extension to SQL.
- TransactSQL (T-SQL is SQL Server's procedural language extension to SQL.
- It provides a server-side, stored procedural language that is easy-to-use, seamless with SQL, robust, portable, and secure.
- PL/SQL enables you to mix SQL statements with procedural constructs, such as loops, conditions, and exceptions.



46

46

```

SET SERVEROUTPUT ON
DECLARE
    CURSOR movie_data IS
        SELECT mvtitle,dirname
            FROM movie, director
            WHERE movie.dirnumb = director.dirnumb;

    director_name director.dirname%type;
    movie_name movie.mvtitle%type;
BEGIN
    OPEN movie_data;
    LOOP
        FETCH movie_data INTO movie_name, director_name;
        EXIT WHEN movie_data%NOTFOUND;

        DBMS_OUTPUT.put_line(movie_name || ' ' ||
director_name);
    END LOOP;
    CLOSE movie_data;
END;
/

```

47

Stored Procedures

- We will have a hands-on PL/SQL discussion next week.

48

Stored Procedures

- An E-Commerce database has two tables: sales and inventory. Whenever a sales transaction is completed, a new sales record is added to sales table and the 'stock_level' value is updated on the corresponding row in the inventory table.
- At the end of the day, inventory table is checked to replenish stocks.
- Should we use a materialized view to check inventory table?
- Should we use a trigger to update inventory table?



49

49

A presentation slide with a dark blue background featuring a white line-art pattern of stylized fish or leaves. On the left, there is a large red semi-circle. On the right, there is a large blue cross shape containing the RMIT University logo. The text 'Menti-time!' is written in large white font, with 'Menti.com' and 'Code: to be supplied' in smaller white font below it. A green horizontal line is above the text, and a blue horizontal line is below it.

50

