

Database Applications

Lecture 8: NoSQL and MongoDB

Santha Sumanasekara
September 2022



1

1

Announcements

- Mid-Semester Test – Thursday 22 Sep --
 - Must bring in your laptop (fully charged)
 - Topics covered in Week 1 – 7 are examinable. No PHP coding
 - Multiple-choice and short-answer questions
- Assignment 2 – will be released next week.




2

2

A (not so) short history of DBMS

- http://upload.wikimedia.org/wikipedia/commons/2/22/RDBMS_timeline.svg



Re-invent the wheel?

- In modern times, there were only very few re-inventions that made a great impact!
 - Propellor engine → Jet engine (in aircraft)
 - Film Camera → Digital Camera
 - ...
- Is NoSQL the database equivalent to the jet engine in the aeronautical world?
- Why do you reinvent this new database model?

Re-invent the wheel?

- Data is everywhere. The digital transformation drives us to the frontiers that never reached.
- The digital transformation demands for rich content and big data!
- The structured database model, that catered the data management needs for almost half a century, has its inherent limitations.



5

5

An Example

- Consider the following example: Consider a database schema developed in early 1990s where use of email was not widespread, however, now it has become an essential piece of information required for a student database.

ID	Name
1	Caleb
2	Sal
3	Ann
4	Jacob
5	Amber

Email does not fit table structure

ann@gmail.com

- This is not easily done with a relational database model.



6

6

Another Example

- Consider the following two AirBnB listings:
 - <https://www.airbnb.com.au/experiences/79594>
 - <https://www.airbnb.com.au/rooms/10082422>
- How different are they?
- Can you easily store them in a table in a relational database?
- We will explore these databases in our future labs.
(<https://docs.atlas.mongodb.com/sample-data/sample-airbnb/#sample-document>)



7

7

NoSQL – Not just one model

<https://en.wikipedia.org/wiki/NoSQL>

NoSQL is a collective noun used for data models other than tabular relations.

Some of the common NoSQL data models are:

- Document databases
- Graph stores
- Key-value stores
- Wide-column stores



8

8

NoSQL Database Types – Document Databases

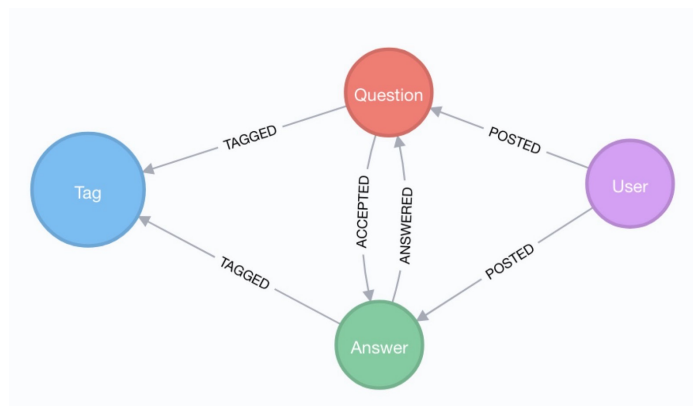
- **Document databases:** pair each key with a complex data structure known as a document. Documents can contain many different key-value pairs, or key-array pairs, or even nested documents.

```
{
  "_id": "10006546",
  "listing_url": "https://www.airbnb.com/rooms/10006546",
  "name": "Ribeira Charming Duplex",
  "summary": "Fantastic duplex apartment with three bedrooms,
    located in the historic area of Porto, Ribeira (Cube)...",
  "interaction": "Cot - 10 € / night Dog - € 7,5 / night",
  "house_rules": "Make the house your home...",
  "property_type": "House",
  "room_type": "Entire home/apt",
  ...
}
```

9

NoSQL Database Types – Graph Stores

- **Graph stores** are used to store information about networks of data, such as social connections. Graph stores include Neo4J and Giraph.



10

NoSQL Database Types – Key-value stores

- **Key-value stores** are the simplest NoSQL databases. Every single item in the database is stored as an attribute name (or 'key'), together with its value. Examples of key-value stores are Riak and Berkeley DB. Some key-value stores, such as Redis, allow each value to have a type, such as 'integer', which adds functionality.



11

11

NoSQL Database Types – Wide-column Stores

- **Wide-column stores** such as Cassandra and HBase are optimized for queries over large datasets, and store columns of data together, instead of rows.



12

12

Benefits of NoSQL

➤ Dynamic Schemas

- Relational databases require that schemas be defined before you can add data.
- NoSQL databases are built to allow the insertion of data without a predefined schema.
- Nicely fits with agile development process.
- Data validation is delegated to the application.



13

13

Benefits of NoSQL

➤ Auto-sharding

- Partitioning can be done in the relational database systems, however, they require quite sophisticated mechanisms to ensure joins etc are executed efficiently.
- NoSQL databases, on the other hand, usually support auto-sharding, meaning that they natively and automatically spread data across an arbitrary number of servers, without requiring the application to even be aware of the composition of the server pool.



14

14

Benefits of NoSQL

➤ Replication

- Most NoSQL databases also support automatic database replication to maintain availability in the event of outages or planned maintenance events.
- Unlike relational databases, NoSQL databases generally have no requirement for separate applications or expensive add-ons to implement replication.



15

15

MongoDB

What a name for a database!

The name of the database was derived from the word *humongous* to support the idea of processing large amount of data.



16

MongoDB

- **MongoDB** is a cross-platform document-oriented database
- MongoDB uses JSON-like documents. A document is a record (sort of like a row in a structured table). For example, if we have a database for users, a document would be one individual user.
- A document is made out of a list of key-value pairs.



```

1  {
2  name: "sue",
3  age: 26,
4  status: "A",
5  Groups: [ "news", "sports" ]
6  }

```

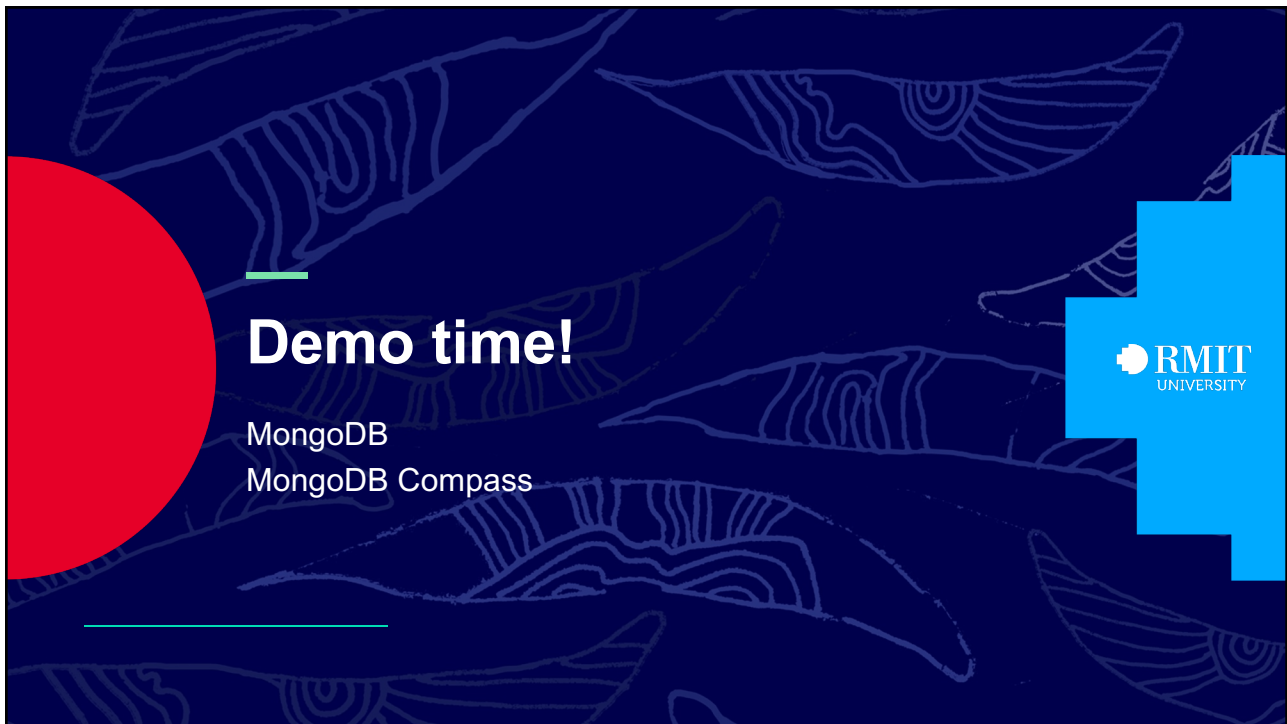
17

MongoDB – similar but different!

Relational	MongoDB
database	database
table	collection
row	document
column	field
<pre>CREATE TABLE people (user_id Varchar(30), age Number, status char(1), PRIMARY KEY (user_id));</pre>	<pre>db.people.insertOne({ user_id: "abc123", age: 55, status: "A" })</pre>

18

18



19

Insert a document

➤ In SQL:

```
INSERT INTO people (user_id, age, status)
VALUES ('A1234', 23, 'A');
```

➤ In MongoDB Shell:

```
db.people.insertOne(
  { user_id: "A1234", age: 23, status: "A" }
)
```

20

Insert a document

➤ In SQL:

```
INSERT INTO people (user_id, age, status)
VALUES ('A1234', 23, 'A');
```

➤ In MongoDB Compass:

Insert Document

1	_id : ObjectId("5d791c0828e5b1700b73486b ")	ObjectId
2	user_id : "A1234 "	String
3	age : 34	Int32
4	status : "A "	String

CANCEL INSERT



21

21

A complex insertion

➤ In SQL: We **cannot** insert composite objects or arrays into tables

```
INSERT INTO people (user_id, name, address, age, sports)
VALUES ('A1234', 'Sam', {1}{Happy St}{Happyville}{3999},
23, ['AFL', 'Cricket', 'Rugby']);
```

➤ In MongoDB Shell:

```
db.people.insertOne(
  { user_id: "A1234", name: "Sam",
    address: { no: 28, street: "Happy St",
               suburb: "Happyville"},
    age: 23, sports: ["AFL", "Cricket", "Rugby"] }
)
```

Address is a composite object. Sports is an array.

22

22

A complex insertion

- In MongoDB Compass:

Insert Document

1	_id : ObjectId("5d798a1e28e5b1700b73...")	ObjectId
2	user_id : "A1234 "	String
3	name : "Same "	String
4	address : Object	Object
5	no : 1	Int32
6	street : "Happy St "	String
7	suburb : "Happyville "	String
8	postcode : 3999	Int32
9	age : 32	Int32
10	sports : Array	Array
11	0 : "AFL "	String
12	1 : "Cricket "	String
13	2 : "Rugby "	String

CANCEL INSERT

Address is a composite object. Sports is an array.



23

23

Find a document

- In SQL we use SELECT to find/ retrieve data. In MongoDB, we use find() method.

```
db.<<collection-name>>.find (<<query>>, <<projection>>)
```

This specifies the conditions.
Similar to WHERE clause in SQL

- Next few slides demonstrate how queries and projections are formed.

This specifies which fields to display. Similar to SELECT clause in SQL.



24

24

Find a document

➤ In SQL:

```
SELECT *
  FROM people;
```

➤ In MongoDB Shell:

```
db.people.find()
```



25

25

Find some fields in a document

➤ In SQL:

```
SELECT ROWID, user_id, status
  FROM people;
```

Leave this empty, as there isn't any filtering condition.

➤ In MongoDB Shell:

```
db.people.find(
  { },
  { user_id: 1, status: 1 }
)
```

If you do not want to see the document ID, add `_id: 0` to the projection (field list).



26

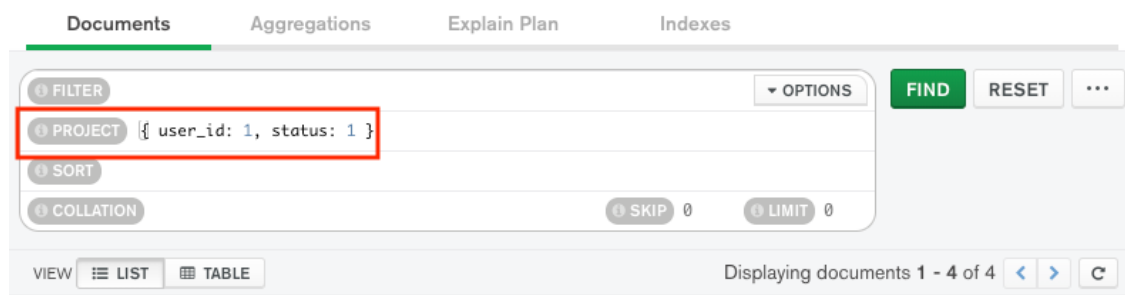
26

Find a document

➤ In SQL:

```
SELECT ROWID, user_id, status
FROM people;
```

➤ In MongoDB Compass:



27

Filter a document

➤ In SQL:

```
SELECT user_id, status
FROM people
WHERE status = 'A';
```

➤ In MongoDB Shell:

```
db.people.find(
  { status: "A" },
  { user_id: 1, status: 1, _id: 0 }
)
```

filtering condition.

28

Filter a document

➤ In SQL:

```
SELECT user_id, status
FROM people
WHERE status = 'A';
```

➤ In MongoDB Compass:

The screenshot shows the MongoDB Compass interface with the 'Documents' tab selected. The 'FILTER' field contains the query `{status: "A"}`. The 'PROJECT' field contains `{ user_id: 1, status: 1 }`. The 'SORT' and 'COLLATION' fields are empty. The 'SKIP' and 'LIMIT' fields are both set to 0. The 'FIND' and 'RESET' buttons are visible on the right.

29

29

Filter a document – with multiple conditions

➤ In SQL:

```
SELECT *
FROM people
WHERE status = 'A' AND
      age = 25;
```

AND is the default.

➤ In MongoDB Shell:

```
db.people.find(
  { status: "A",
    age: 25 }
)
```

30

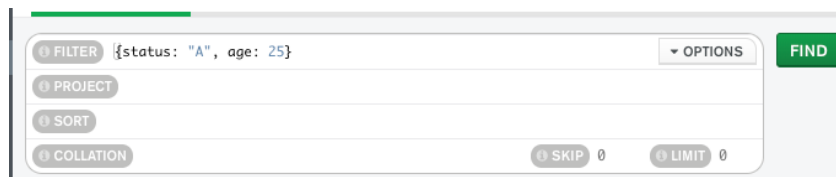
30

Filter a document

➤ In SQL:

```
SELECT *
  FROM people
 WHERE status = 'A' AND
        age = 25;
```

➤ In MongoDB Compass:



The screenshot shows the MongoDB Compass interface. The 'FILTER' tab is active, displaying the query `{status: "A", age: 25}`. Below the filter, there are tabs for 'PROJECT', 'SORT', and 'COLLATION'. At the bottom, there are 'SKIP' and 'LIMIT' options, both set to 0. A green 'FIND' button is located on the right side of the interface.



31

31

Filter a document – with OR conditions

➤ In SQL:

```
SELECT *
  FROM people
 WHERE status = 'A' OR
        age = 25;
```

OR is applied to an array of conditions. [] denotes an array.

➤ In MongoDB Shell:

```
db.people.find(
  { $or: [ { status: "A" } , { age: 25 } ] }
)
```



32

32

Filter a document – with OR conditions

➤ In SQL:

```
SELECT *
  FROM people
 WHERE status = 'A' OR
        age = 25;
```

➤ In MongoDB Compass:



The screenshot shows the MongoDB Compass interface. The 'FILTER' tab is active, displaying the query: `{ $or: [{ status: "B" }, { age: 25 }] }`. Below the filter, there are tabs for 'PROJECT', 'SORT', and 'COLLATION'. On the right, there are 'FIND' and 'RESET' buttons. At the bottom right, there are 'SKIP' and 'LIMIT' buttons, both set to 0.



33

33

Filter a document – with “Not Equal”

➤ In SQL:

```
SELECT *
  FROM people
 WHERE age <> 25;
```

\$ne means “not equal”

➤ In MongoDB Shell:

```
db.people.find(
  { age: { $ne: 25 } }
)
```



34

34

Filter a document – middle-aged people!

➤ In SQL:

```
SELECT *
  FROM people
 WHERE age > 35 AND age <= 60;
```

➤ In MongoDB Shell:

```
db.people.find(
  { age: { $gt: 35, $lte: 60 } }
)
```

\$gt: greater than
\$gte: greater than or equal
\$lt: less than
\$lte: less than or equal



35

35

Filter a document – Partial Matches

➤ In SQL:

```
SELECT *
  FROM people
 WHERE name LIKE '%Sam%';
```

➤ In MongoDB Shell:

```
db.people.find(
  { name: /Sam/ }
)
```

More complex regular expressions are possible. To be discussed later. E.g.
{ name: { \$regex: /Sam/ } }



36

36

Filter a document – Using Composite Fields

- In SQL: Cannot be done!

```
SELECT *  
  FROM people  
 WHERE address.suburb = 'Happyville';
```

- In MongoDB Shell:

```
db.people.find(  
  {"address.suburb": "Happyville"}  
)
```



Make sure to use double quotes.

37

37

Filter a document – Using Array Values

- In SQL: Cannot be done!

```
SELECT *  
  FROM people  
 WHERE sport = ['AFL', 'Cricket']
```

- In MongoDB Shell:

```
db.people.find(  
  {sports:"AFL", "Cricket"}  
)
```



Return documents that exactly contain this array.

38

38

Aggregations -- count()

➤ In SQL:

```
SELECT COUNT(*)
  FROM people;
```

➤ In MongoDB Shell:

```
db.people.find().count()
```

Object-oriented. Two ways to do it:
`db.collection.find().count()`
`db.collection.count()`



39

39

Sorting

➤ In SQL:

```
SELECT *
  FROM people
 ORDER BY name;
```

➤ In MongoDB Shell:

```
db.people.find().sort({name: 1})
```

```
db.people.find().sort({name: -1})
```



Descending order

40


40

Sorting

➤ In SQL:

```
SELECT *  
  FROM people  
  ORDER BY name;
```

➤ In MongoDB Compass:



The screenshot shows the MongoDB Compass interface. The 'SORT' tab is selected, and the query field contains '{name: 1}'. The 'SKIP' and 'LIMIT' fields are both set to 0. A 'FIND' button is visible on the right.

Further Readings

- <https://docs.mongodb.com/manual/reference/sql-comparison/>
- <https://docs.mongodb.com/manual/tutorial/insert-documents/>
- <https://docs.mongodb.com/manual/tutorial/query-documents/>

Next Week

- Mid-semester Test
- We will continue MongoDB discussion in Week 9. We discuss further on aggregations, etc.



43

43

Menti-time!

Menti.com

Code: to be supplied



44

